

Are my results reproducible?

PARALLEL PROGRAMMING IN R



Hana Sevcikova

University of Washington

Random numbers in R

- Many statistical applications involve random numbers (RNs)
- Examples: MCMCs in Bayesian methods, bootstrap, simulations

```
set.seed(1234)  
rnorm(3)
```

```
-1.2070657  0.2774292  1.0844412
```

```
rnorm(3)
```

```
-2.3456977  0.4291247  0.5060559
```

- For reproducibility:
 - Set **seed of a random number generator** (RNG) prior to running the code

```
set.seed(1234)  
rnorm(3)
```

```
-1.2070657  0.2774292  1.0844412
```

```
rnorm(3)
```

```
-2.3456977  0.4291247  0.5060559
```

Naive (non)reproducibility in parallel code

```
library(parallel)
cl <- makeCluster(2)
set.seed(1234)
clusterApply(cl, rep(3, 2), rnorm)
```

```
[[1]]
[1] -1.891091 -1.351767 -1.456848

[[2]]
[1]  1.7346577  0.7855641 -2.2319774
```

```
set.seed(1234)
clusterApply(cl, rep(3, 2), rnorm)
```

```
[[1]]
[1]  0.4432499 -0.7896067  0.2659675

[[2]]
[1]  0.2229560  0.8323269 -0.4092570
```

Incorrect way of generating RNs in parallel code

- Using `set.seed()`, the RNG is initialized only on the master.
- Workers start with a clean environment, thus no RNG seed set.
- What happens when we set the RNG on each worker?

```
clusterEvalQ(cl, set.seed(1234))  
clusterApply(cl, rep(3, 2), rnorm)
```

```
[[1]]  
[1] -1.2070657  0.2774292  1.0844412  
  
[[2]]  
[1] -1.2070657  0.2774292  1.0844412
```

- Quick and dirty solution:

```
for (i in 1:2) {  
  set.seed(1234)  
  clusterApply(cl, sample(1:10000000, 2), set.seed)  
  print(clusterApply(cl, rep(3, 2), rnorm))}
```

```
[[1]]  
[1] 0.078249533 0.003019703 -1.314239709  
[[2]]  
[1] 1.3955357 -0.9935141 -0.3740712  
[[1]]  
[1] 0.078249533 0.003019703 -1.314239709  
[[2]]  
[1] 1.3955357 -0.9935141 -0.3740712
```

- NOT RECOMMENDED!!!

Let's practice!
PARALLEL PROGRAMMING IN R

Parallel random number generators

PARALLEL PROGRAMMING IN R



Hana Sevcikova
University of Washington

Random Number Generators (RNGs)

- Important parameters of an RNG:
 - long period (preferably $> 2^{100}$)
 - good structural (distributional) properties in high dimensions
- These parameters should hold when used in distributed environment

L'Ecuyer Multiple Streams RNG

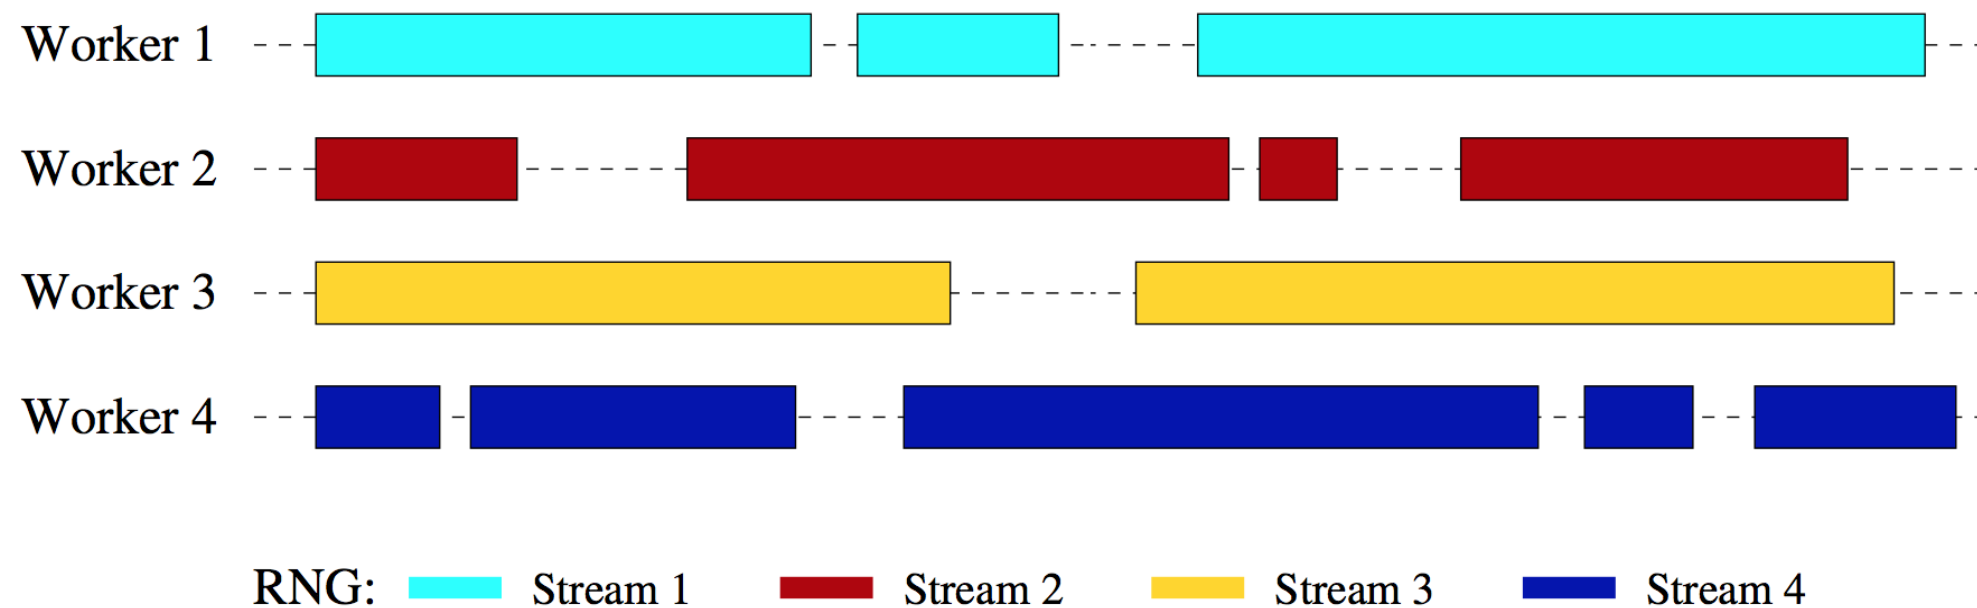
- A good quality RNG with multiple independent streams proposed by Pierre L'Ecuyer et al. (2002), RngStreams
 - Period 2^{191}
 - Streams have seeds 2^{127} steps apart
 - Parallel parts of user computation can use independent and reproducible streams
 - Direct interface in R: `rlecuyer` , `rstream`
 - In R core: `RNGkind("L'Ecuyer-CMRG")`

Using L'Ecuyer RNG in parallel

- Setting an RNG seed for cluster `cl`:

```
clusterSetRNGStream(cl, iseed = 1234)
```

- Initializes a reproducible independent stream on each worker



Reproducibility in the parallel package

- In `parallel` : one stream per worker
- Creates constraints on reproducibility
- Results only reproducible if:
 1. process runs on clusters of the same size
 2. process does **not** use load balancing, e.g. `clusterApplyLB()`

Let's practice!
PARALLEL PROGRAMMING IN R

Reproducibility in foreach and future.apply

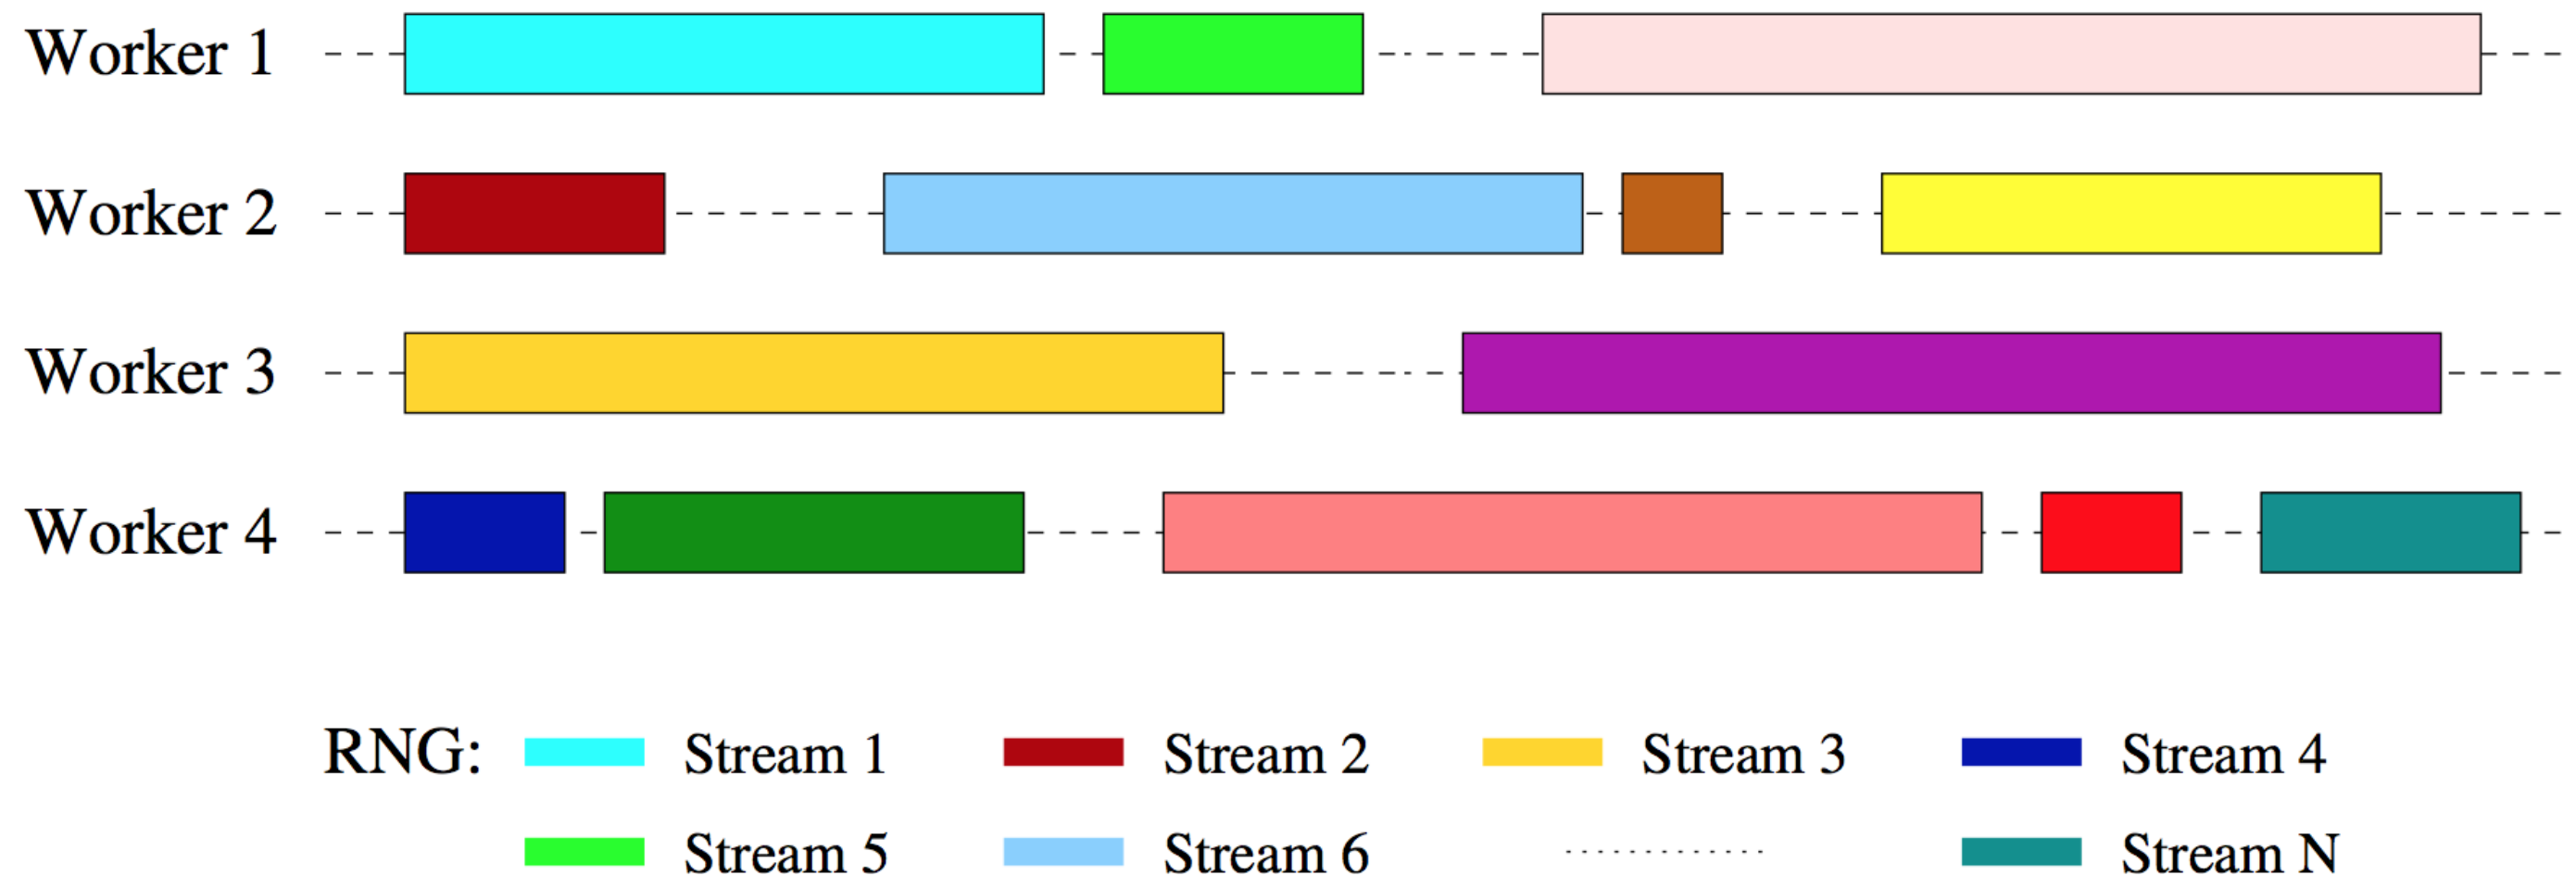
PARALLEL PROGRAMMING IN R



Hana Sevcikova
University of Washington

doRNG: backend for foreach

- Reproducibility in `foreach` via R package `doRNG` :
 - Uses independent streams of the L'Ecuyer's RNG
 - Instead of one stream per worker, it uses **one stream per task**



Using doRNG via %dorng%

```
library(doRNG)
library(doParallel)
registerDoParallel(cores = 3)
```

```
set.seed(1)
res1 <- foreach(n = rep(2, 5), .combine = rbind) %dorng% rnorm(n)
```

```
set.seed(1)
res2 <- foreach(n = rep(2, 5), .combine = rbind) %dorng% rnorm(n)
```

```
identical(res1, res2)
```

```
TRUE
```

Using doRNG via %dopar%

```
library(doRNG)
library(doParallel)
registerDoParallel(cores = 3)
```

```
registerDoRNG(1)
res3 <- foreach(n = rep(2, 5), .combine = rbind) %dopar% rnorm(n)
```

```
set.seed(1)
res4 <- foreach(n = rep(2, 5), .combine = rbind) %dopar% rnorm(n)
```

```
c(identical(res1, res3), identical(res2, res4))
```

```
TRUE TRUE
```

Note: Cannot be used with the `%doSEQ%` backend.

Summary of using doRNG

Two ways of including `doRNG` into `foreach` :

1. Using `%dorng%` :
 - advantage of being explicit about using the L'Ecuyer's RNG
2. Using `%dopar%` and registering `doRNG` :
 - easy to make code/packages reproducible by only prepending `registerDoRNG()`

`doRNG` can be used with any parallel backend, including `doFuture` .

future.apply

- Uses independent streams of the L'Ecuyer's RNG
- As in `doRNG` , generates one stream per task
- Need only to assign `future.seed` argument

```
library(future.apply)
plan(sequential)
res5 <- future_lapply(1:5, FUN = rnorm, future.seed = 1234)
plan(multiprocess)
res6 <- future_lapply(1:5, FUN = rnorm, future.seed = 1234)
identical(res5, res6)
```

TRUE

Let's practice!
PARALLEL PROGRAMMING IN R

Finishing Touch

PARALLEL PROGRAMMING IN R



Hana Sevcikova

Senior Research Scientist, University of
Washington

Recommended R packages (1)

- `parallel` (core package)
 - No need for dependencies on other packages
 - Important to understand as other packages are built on it
 - Often yields best performance
 - Reproducible results: only on clusters of the same size with no load balancing

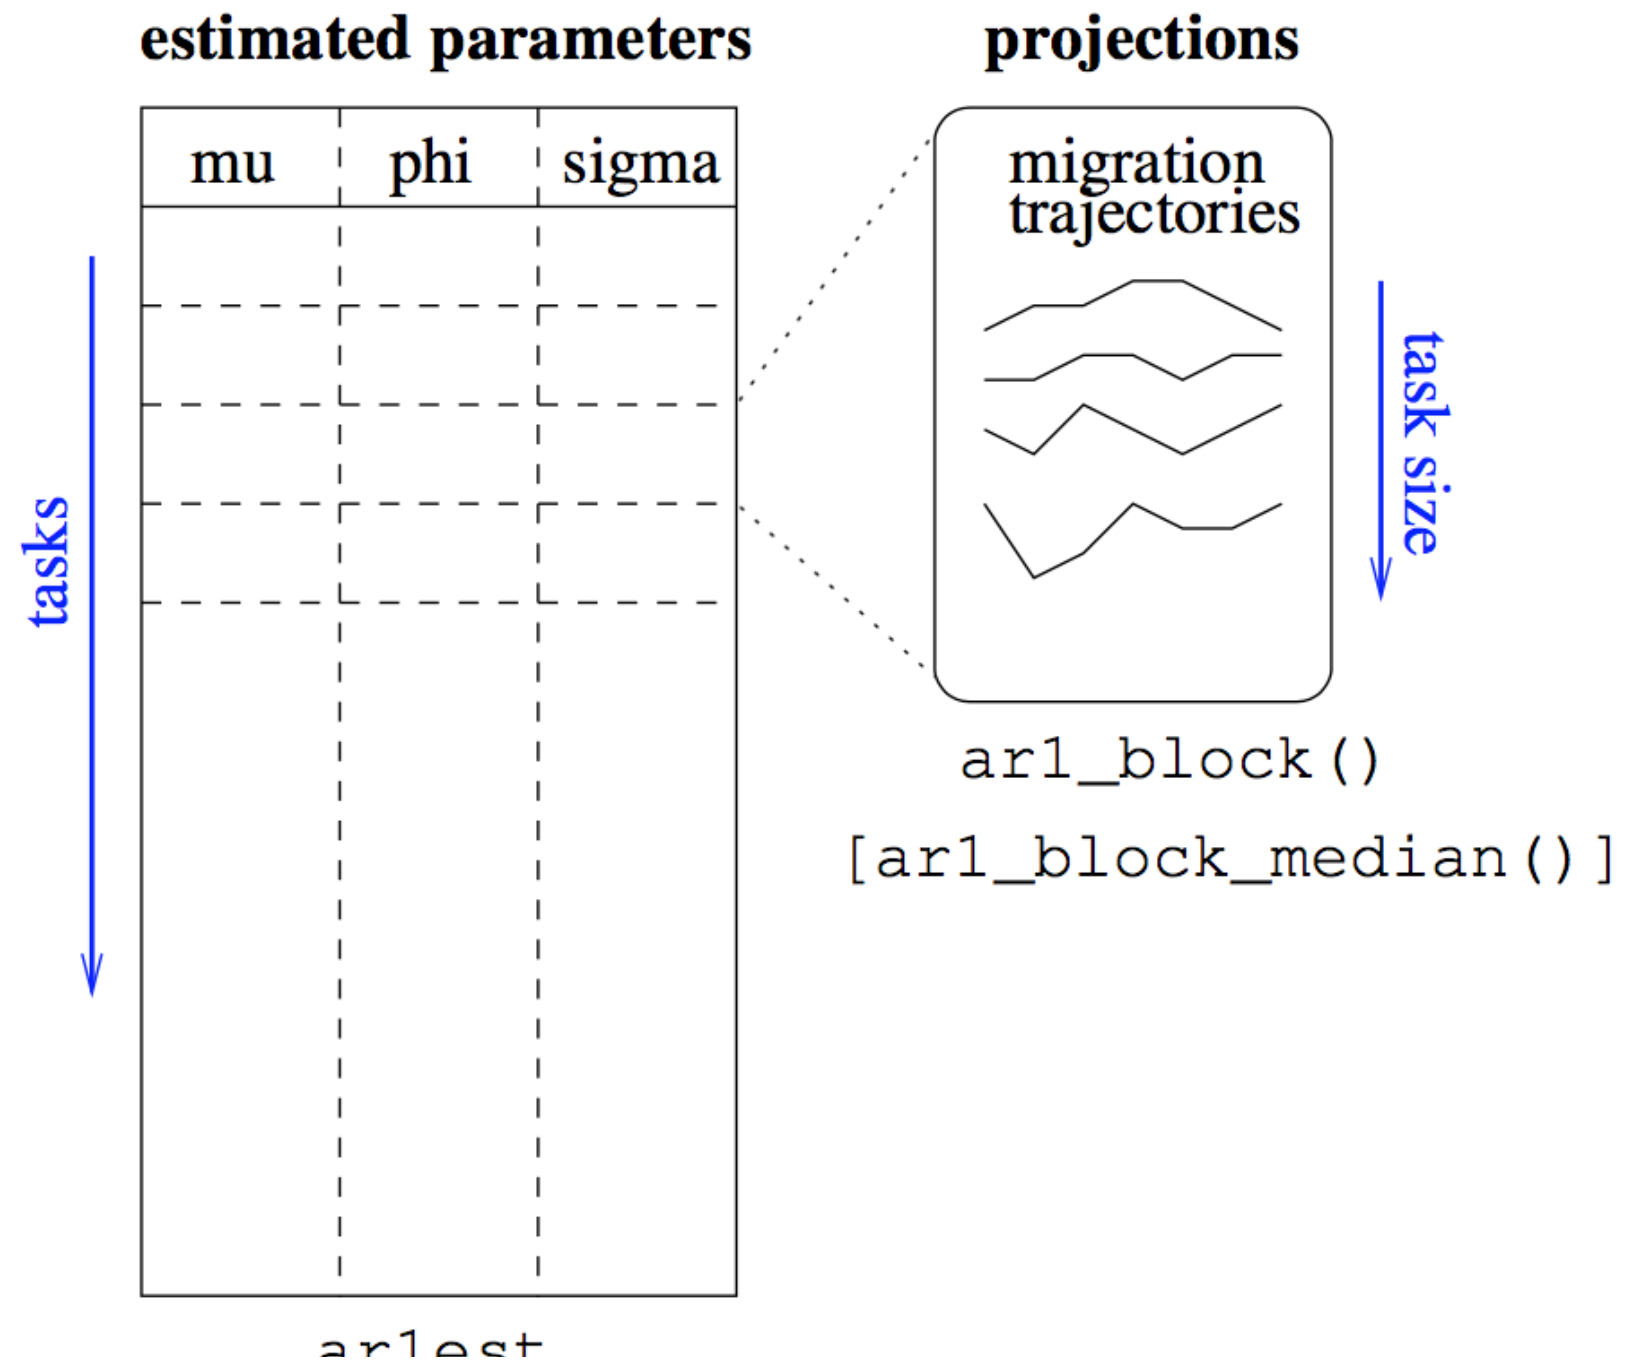
Recommended R packages (2)

- **foreach** (with **doParallel**, **doFuture**)
 - Higher level programming
 - Intuitive syntax in form of **for** loops
 - Results reproducible via **doRNG**
- **future.apply** (based on **future**)
 - Unifies many parallel backends into one interface
 - Intuitive **apply()**-like syntax
 - Results always reproducible


Getting the best performance

- Minimize amount of **communication** (sending repeatedly big data is bad!)
- Use **scheduling** and **load balancing** appropriate for your application (e.g. group tasks into chunks evenly distributed across workers)
- Use **cluster size** appropriate for your hardware (i.e. number of physical cores)

Probabilistic projection of migration



Speedup: $T_{\text{sequential}}/T_{\text{parallel}}$



Congratulations!

PARALLEL PROGRAMMING IN R